

# SECTION SECTION 3 Dimension Reduction

## Curse of Dimensionality (COD)

Suppose that we have a one variable data set with 1,000,000 cases. If we want a histogram with our variable divided into 10 intervals (bins) we would have, on average, 100,000 cases/bin.

Now suppose the data is in two dimensions. If we want to have each variable divided into 10 intervals we would have 100 bins and, on average, 10,000 cases/bin.

Dim	Bins	Ave. Cases/bin
1	10	100,000
2	100	10,000
3	1,000	1,000
4	10,000	100
5	100,000	10
6	1,000,000	1
7	10,000,000	0.1
$\vdots$	$\vdots$	$\vdots$
10	10,000,000,000	0.0001

We see that the data is sparsely distributed in the bins.

It is also interesting to consider the nature of the hypercubes. If we have a box in 3 dimensions

$[(1, 1, 1), (1, 1, -1), (1, -1, -1), (1, -1, 1), (-1, 1, 1), (-1, 1, -1), (-1, -1, -1), (-1, -1, 1)]$ , the distance from  $(0, 0, 0)$  to any corner is  $\sqrt{3}$ , while in 4 dimensions it is  $\sqrt{4}$ , etc. As a result, the distance from the centre to a corner is increasing, while the distance from the centre to an edge is constant (1) so a greater proportion of the volume will be in the corners. Almost every point is closer to an edge than to another point. This sparseness problem is commonly called the "curse of dimensionality" (COD).

Since we use samples to estimate an unknown function, our estimates may be inaccurate (biased). Meaningful estimation is possible only for sufficiently *smooth* functions but sparseness of high-dimensional space makes it difficult to collect enough samples to attain a high enough density to ensure a sufficiently smooth function. Smoothness constraints describe how individual cases in the training data are combined by the learning method in order to construct the function estimate. Accurate function estimation depends on having

## Data Science

enough cases within the neighbourhood specified by the smoothness constraints. As the number of dimensions increases, the number of cases needed to give the same density increases exponentially. This could be offset by increasing the neighbourhood size with dimensionality (increasing the number of cases falling within the neighbourhood) but this is at the expense of imposing stronger (and possibly incorrect) constraints. Low data density requires us to specify stronger, more accurate constraints on the problem solution.

The COD is due to the geometry of high-dimensional spaces. A large radius is needed to enclose a fraction of the data points in a high-dimensional space (see above example). For a given fraction of cases, it is possible to determine the edge length of the hypercube using

$$e_d(p) = p^{1/d}$$

where  $p$  is the (prespecified) fraction of cases. In a 10-dimensional space ( $d = 10$ ) to enclose 10% of the cases, the edge length is  $e_{10}(0.1) = 0.80$ . Thus very large neighbourhoods are needed to capture even a small portion of the data. For a sample of size  $n$ , the expected  $L_\infty$  distance between data points is

$$D(d, n) = 1/2(1/n)^{1/d}$$

so for a 10-dimensional space,  $D(10, 1000) \approx 0.5$  and  $D(10, 10,000) \approx 0.4$ .

Thus there are serious problems associated with making local estimates for high-dimensional samples and a lot of extrapolation will be required.


## 4.1 Reducing Dimensionality

---

In data that has very high dimensions, it can be important to reduce the effective dimension of the data to enable us to employ some methods that work better in lower dimensions. For example, doing an All Subsets Regression on 12 variables will be far easier than on 125 variables.

## An example of dimension reduction


Consider data points at  $(1, 1)$ ,  $(2, 2)$ ,  $(3, 3)$ ,  $(4, 4)$ ,  $(5, 5)$ .



**Figure 1.**

These points are specified in terms of the two orthonormal vectors  $\mathbf{e}_1 = (1, 0)$  and  $\mathbf{e}_2 = (0, 1)$ .

What happens if we instead use the two orthonormal vectors  $\mathbf{e}'_1 = (1/\sqrt{2}, 1/\sqrt{2})$  and  $\mathbf{e}'_2 = (-1/\sqrt{2}, 1/\sqrt{2})$ ? This gives us a new set of axes as shown.



**Figure 2.**

Because the points all lie on the basis vector  $\mathbf{e}'_1$  we can ignore the other basis vector for our new coordinate system. This results in a reduction of the dimension of our dataset.

What happens if some or all of the data points are not exactly on the line? If they are not too far off, we may feel that we can ignore the slight difference and represent the data in terms of just one vector.

## 4.2 Principal Components Analysis (PCA)

---

PCA is one of the standard methods for dimension reduction.

When one looks at data with a scatterplot matrix, the only structure that can be seen is that which is visible from the original coordinate axes and is restricted to relationships between two variables. If we look at a scatterplot matrix and find that every scatterplot that involved one of the variables is virtually a horizontal or vertical straight line then we would conclude that we could model the data without using that variable. We might find, if we rotate the data, that a similar behaviour could be seen for **linear combinations of the variables**. PCA is a method that enables us to see if such structure exists.

Consider our data as random variables  $X_1, \dots, X_p$  with  $n$  observations for each of these random variables. Principal components are special linear combinations of the  $p$  random variables. These linear combinations represent a *new coordinate system* that is obtained by rotating the original system that had  $X_1, \dots, X_p$  as the coordinate axes. The new axes represent directions of variability and provide a simpler, more parsimonious description of the covariance structure of the original data. Principal components depend solely on the covariance matrix  $\Sigma$  (or the correlation matrix  $\rho$ ) of  $X_1, \dots, X_p$ .

Specifically, *PCA* looks at variance in the data and identifies the mutually orthogonal directions of decreasing variance. In *PCA* we form as many new variables as we have original variables. The new variables are linear combinations of the old variables. But they are chosen in such a way that the first linear combination (*PC1*) explains the highest proportion of the variance in the original variables. The second linear combination (*PC2*) is orthogonal to the first and it explains the second largest proportion of variance in the original variables. The third linear combination (*PC3*) is chosen to be orthogonal to the first two and it explains the third largest proportion of the variance in the original variables, etc.

Let the random vector  $\mathbf{X}^T = [X_1, \dots, X_p]$  have covariance matrix

$$\Sigma = \begin{bmatrix} \text{var}(X_1) & \text{cov}(X_1, X_2) & \cdots & \text{cov}(X_1, X_p) \\ \text{cov}(X_2, X_1) & \text{var}(X_2) & \cdots & \text{cov}(X_2, X_p) \\ \vdots & \cdots & \ddots & \vdots \\ \text{cov}(X_p, X_1) & \cdots & \cdots & \text{var}(X_p) \end{bmatrix}$$

with eigenvalues  $\lambda_1 \geq \dots \geq \lambda_p \geq 0$  and associated eigenvectors. We shall see that eigenvalues and eigenvectors have specific statistical interpretations. Consider  $p$  linear combinations of the original  $p$  variables, i.e.

$$Y_i = \ell_i^T \mathbf{X}, i = 1, \dots, p$$

Then

$$\text{var}(Y_i) = \ell_i^T \Sigma \ell_i$$

and

$$\text{cov}(Y_i, Y_k) = \ell_i^T \sum \ell_k \quad \text{for } i, k = 1, \dots, p$$

To eliminate indeterminacy, we restrict ourselves to coefficient vectors of length one. Hence  $PC1$  (i.e.  $Y_1$ ) is the linear combination  $\ell_1^T \mathbf{X}$  that maximizes  $\text{var}(\ell_1^T \mathbf{X})$  subject to

$$\ell_1^T \ell_1 = 1;$$

$PC2$  (i.e.  $Y_2$ ) is the linear combination  $\ell_2^T \mathbf{X}$  that maximizes  $\text{var}(\ell_2^T \mathbf{X})$  subject to

$$\ell_2^T \ell_2 = 1$$

and

$$\text{cov}(Y_1, Y_2) = 0$$

etc.

We can show that

$$\max_{\ell \neq 0} \frac{\ell^T \sum \ell}{\ell^T \ell} = \lambda_1 = \text{Var}(Y_1)$$

and is attained when  $\ell = \mathbf{e}_1$  and

$$\max_{\ell \perp \mathbf{e}_1, \dots, \mathbf{e}_k} \frac{\ell^T \sum \ell}{\ell^T \ell} = \lambda_{k+1} (= \text{Var}(Y_{k+1})) \quad \text{for } k = 1, 2, \dots, p-1$$

and

$$\sum_{i=1}^p \text{var}(X_i) = \sum_{i=1}^p \text{var}(Y_i).$$

We have thus constructed independent components that are conserving the total variance in the dataset but we have not yet reduced the dimension. To reduce dimension we may drop the latter  $PC$ s which explain less of the variance in the original data.

## Data Science

We can illustrate on the simple example (Figure 1) described above.

```
library(stats)
A <- cbind(1:5,1:5)
(A.pc <- prcomp(A))
Standard deviations:
[1] 2.236068e+00 1.431424e-16
Rotation:
      PC1      PC2
[1,] 0.7071068 -0.7071068
[2,] 0.7071068  0.7071068
summary(A.pc)
Importance of components:
      PC1 PC2
Standard deviation      2.24 1.43e-16
Proportion of Variance 1.00 0.00e+00
Cumulative Proportion 1.00 1.00e+00
(A.new <- A%%A.pc$rotation[,1])
      [,1]
[1,] 1.414214
[2,] 2.828427
[3,] 4.242641
[4,] 5.656854
[5,] 7.071068
A.new%%A.pc$rotation[,1]
      [,1] [,2]
[1,]      1      1
[2,]      2      2
[3,]      3      3
[4,]      4      4
[5,]      5      5
```

We can see that principal components produce the transformation that we expected and allow us to obtain the coordinates in the new coordinate system that corresponds to the eigenvectors (note  $PC1$  is  $\mathbf{e}_1'$  and  $PC2$  is  $\mathbf{e}_2'$ ), as well as the process for returning to the original coordinate system.

Now consider a more realistic problem.

Consider the following 2-dimensional ellipse, with 500 points from a random uniform distribution:

```
numb<- 500          # data set with 500 points
a <- 10              # semi-major axis
b <- 5               # semi-minor axis
x <- runif(numb,-a,a) # x is random uniform in [-a, a] i.e. U[-10,10]
y <- matrix(0,1,numb)
for (i in (1:numb)) {
  aa <- b*(1 - (x[i]/a)^2)^(1/2)
  y[i] <- runif(1,-aa,aa) # a random number in U[-aa, aa]
}
y <- as.vector(y)
plot(x, y, pch = 20, main = "Ellipse? - default scaling", cex = 1.5)
# To make it look like an ellipse.
plot(x,y, pch = 20, asp = 1, main = "Ellipse with correct scaling", cex = 1.5)
```

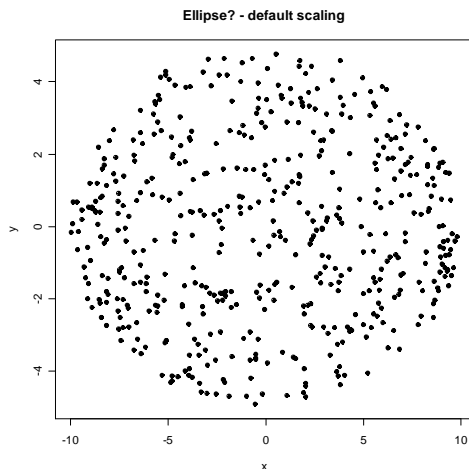


Figure 3.

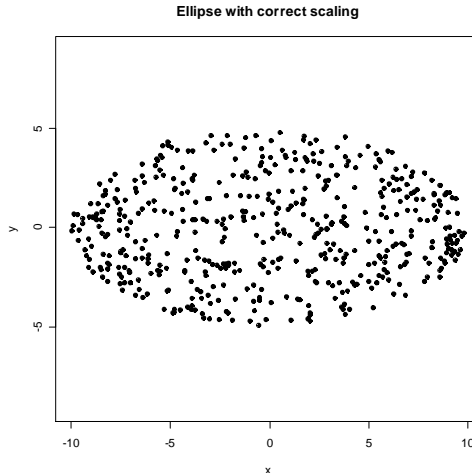


Figure 4.

**Note:** One concern in plotting data is ensuring that the **scaling** on the axes is correct. The left hand graph is the *default* plotting of an elliptic cloud. It looks like a circle because the usual default on plotting is to fill the graphic window as much as possible. The right hand graph is a better representation because it has the same scaling units on both axes.

It is possible that a figure window may be covered by other windows. When that happens, changes to the figure will be done but not be seen. The use of the command

```
bringToTop(which = dev.cur())
```

can assist by either displaying the figure or flashing it on the taskbar.

Now set up the directories

```
drive <- "D:"
code.dir <- paste(drive, "DataMining/Code", sep = "/")
data.dir <- paste(drive, "DataMining/Data", sep = "/")
```

Suppose we now rotate our elliptic cloud by  $\pi/3$  and display it with the original.

```
(R <- cbind(c(cos(pi/3), sin(pi/3)), c(sin(pi/3), -cos(pi/3))))
      [,1] [,2]
[1,] 0.5000000 0.8660254
[2,] 0.8660254 -0.5000000
Z <- cbind(x,y)
XX <- Z%*%R
```

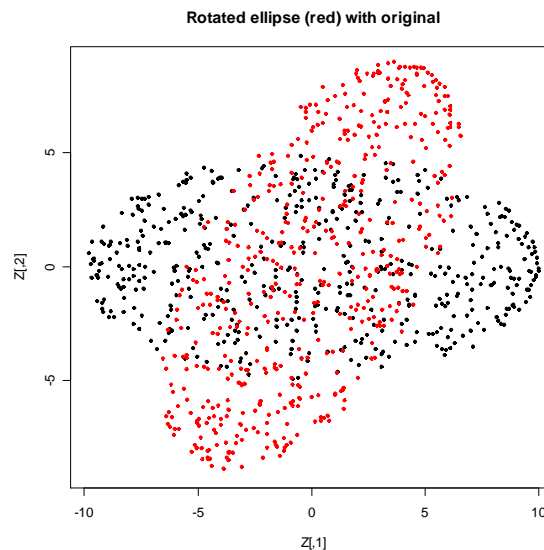
Save the data the first time and read it in the rest of the time so that we can replicate results.

```
# write.table(XX, row.names = F, col.names = F, quote = F, file = paste(data.dir,
"PC_XX.dat", sep = ""))
XX <- read.table(paste(data.dir, "PC_XX.dat", sep = "/"))
```

## Data Science

We need to convert the data frame to a matrix to allow multiplication

```
Z <- as.matrix(XX)%*%R
plot(Z, pch = 20, asp = 1, main = "Rotated ellipse (red) with original")
points(XX, pch = 20, col = "red")
```



**Figure 5.**

It is obvious that the larger variance was along the  $x$ -axis; now it will be along a line at an angle of  $\pi/3$ .

We can find the principal components of this data using the Singular Value Decomposition (SVD) of the covariance matrix  $\Sigma$  of the data.

```
(pc.1 <- svd(cov(XX)))
sqrt(pc.1$d)
```

The Singular Value Decomposition is

```
$d
[1] 33.306878 5.528184
$u
      [,1]      [,2]
[1,] -0.4774353 -0.8786669
[2,] -0.8786669  0.4774353
$v
      [,1]      [,2]
[1,] -0.4774353 -0.8786669
[2,] -0.8786669  0.4774353
```

The standard deviations are

```
[1] 5.771211 2.351209
```

Alternatively, using the `eigen` command, we can obtain the principal components of the



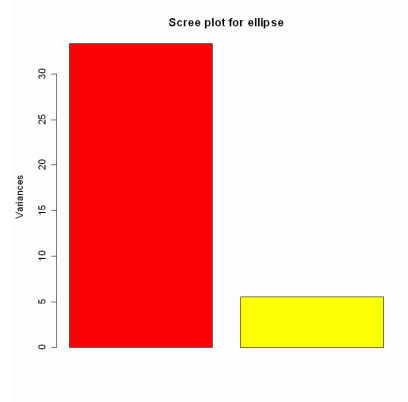
‘new’ data set.

```
(pc.2 <- eigen(cov(XX)))
$values
[1] 33.306878 5.528184
$vectors
      [,1]      [,2]
[1,] -0.4774353 -0.8786669
[2,] -0.8786669  0.4774353
```

(We might get different signs because the direction along the vector is not specified.)

We can also use the principal components method that is given in the library `stats`. It computes other quantities as well.

```
(pc.3 <- prcomp(XX))
Standard deviations:
[1] 5.771211 2.351209
Rotation:
PC1 PC2
V1 -0.4774353 -0.8786669
V2 -0.8786669  0.4774353
plot(pc.3, main = "Scree plot for ellipse", col = c("red", "yellow"))
```



**Figure 6.**

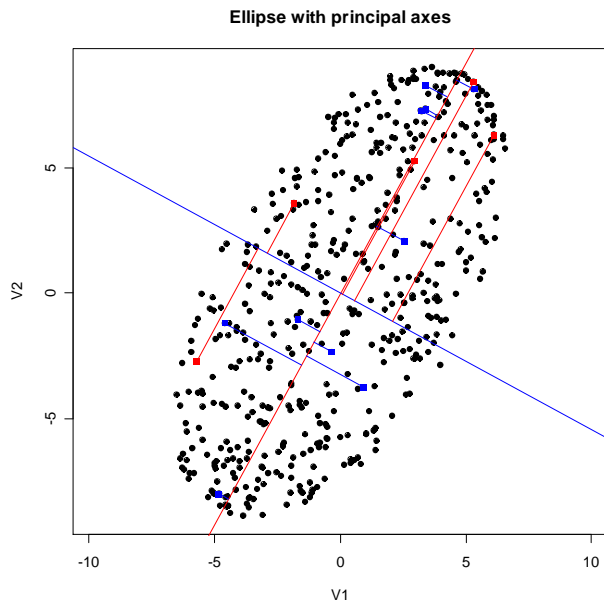
The *scree* plot gives an idea of the relative importance of the principal components since it plots the variance (i.e. the eigenvalues) explained by each successive principal component. Note that they will necessarily be in decreasing order.

```
summary(pc.3)
Importance of components:
              PC1    PC2
Standard deviation  5.771 2.351
Proportion of Variance 0.858 0.142
Cumulative Proportion 0.858 1.000
```

## Data Science


Now we can plot the principal axes (i.e. the eigenvectors) on the data. (`abline` plots lines using the slope and intercept.)

```
plot(XX, pch = 20, asp = 1, main = "Ellipse with principal axes", cex = 1.5)
A <- diag(c(1, 1))
PCaxes <- A%%pc.3$rotation
abline(0,pc.3$rotation[2,1]/pc.3$rotation[1,1], col = "red")
abline(0,pc.3$rotation[2,2]/pc.3$rotation[1,2], col = "blue")
m.1 <- pc.3$rotation[2,1]/pc.3$rotation[1,1]
m.2 <- pc.3$rotation[2,2]/pc.3$rotation[1,2]
for (i in sample(1:500,10)) {
  x.1 <- XX[i,1]
  y.1 <- XX[i,2]
  points(x.1, y.1, pch = 15,col = "blue")
  x.2 <- (y.1 - m.2*x.1)/(m.1 - m.2)
  y.2 <- m.1*x.2
  lines(c(x.1,x.2), c(y.1,y.2), col = "blue")
}
for (i in sample(1:500,5)) {
  x.1 <- XX[i,1]
  y.1 <- XX[i,2]
  points(x.1, y.1, pch = 15,col = "red")
  x.2 <- (y.1 - m.1*x.1)/(m.2 - m.1)
  y.2 <- m.2*x.2
  lines(c(x.1,x.2), c(y.1,y.2), col = "red")
}
```



**Figure 7.** The principal axes are plotted on the cloud (red for the 1st component).

The variance is computed using the sum of the squares of the perpendicular distances from the point to the component axes.



**Figure 8.** The red cloud is the cloud after we have used the principal component rotation to make the 1st principal component axis horizontal.

```
XXR <- as.matrix(XX)%*%PCaxes
plot(XX, pch = 20, asp = 1, main = "Ellipse rotated by PCA rotation", cex = 1.5)
points(XXR, col = "red", pch = "+", cex = 1.5)
```

We need to use `as.matrix(XX)` to convert the data frame values to a matrix.

To see how well PCA determined the orientation of the ellipse, we will compare our known rotation that we applied to the data with the rotation obtained from PCA.

The original ellipse was rotated using the matrix

```
R
      [,1]      [,2]
[ 1,] 0.5000000 0.8660254
[ 2,] 0.8660254 -0.5000000
```

**PCaxes**

The Principal Components rotation is

```
      PC1      PC2
[ 1,] -0.4774353 -0.8786669
[ 2,] -0.8786669  0.4774353
```

## Data Science

It is possible that the first principal component ( $PC1$ ) could be an adequate summary of the total variance in the data provided we feel that the deviations from that line are within our acceptable level of error. (i.e. Does the first Principal Component  $PC1$  explain a sufficiently high proportion of the total variance?) In that case we would have reduced the dimension from two to one.

One difficulty is that by computing principal component variables  $PC1$ ,  $PC2$ , etc., we are computing linear combinations of the original variables so we have probably moved away from physical variables (which have an interpretation) to linear combinations of variables which may have no interpretation. Note that we have not reduced the dimension of our data - we have as many new principal component variables as we had original variables- but if we are willing to **discard** some of the new principal component variables because they account for a very small proportion of the total variance in the dataset, we can reduce the dimension of our problem. We may then be able to use methods that apply to lower dimensional data. Keep in mind that we will then be using the first few principal components (which are linear combinations of the original variables) so we may lose some ability to interpret results.

### Consider the flea beetle example.

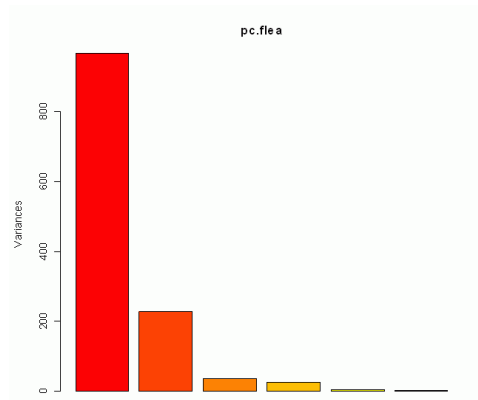
Read in files required for a flea beetle example.

```
source(paste(code.dir, "ReadFleas.r", sep = "/"))
```

PCA gives-

```
(pc.flea <- prcomp(d.flea))
Standard deviations:
[ 1] 31.108528 15.053393  5.976129  5.079179  1.973860  1.095499
Rotation:
      PC1      PC2      PC3      PC4      PC5      PC6
tars1 -0.93222707 -0.32027041  0.14174096  0.06777484  0.02191107 -0.056705703
tars2  0.02339181 -0.43942897 -0.87177026  0.12101445 -0.17802317 -0.005905547
head   0.01753944 -0.12027844 -0.13020947 -0.02687343  0.90584022  0.383408753
aede1  0.15580879 -0.53893004  0.17396789 -0.80548829 -0.04659334 -0.063489191
aede2 -0.05351583 -0.01036395  0.06991407 -0.03878271 -0.38068526  0.919627805
aede3  0.32087002 -0.63190912  0.40965694  0.57421634 -0.01374570 -0.001067120
```

```
plot(pc.flea, col = heat.colors(6))
```



**Figure 9.**

We see that the first principal component is quite dominant and is

$$PC1 = -0.93 tars1 + 0.023 tars2 + 0.018 head + 0.16 aede1 - 0.054 aede2 + 0.32 aede3$$

(Recall that the importance of *tars1* and *aede3* was seen in Ggobi.)

### Consider an example that has 256 variables.

For automated mail sorting, the U. S. Postal Service needs to be able to convert Zip Codes (machine-produced or handwritten) into the corresponding digits. [http://www-stat.stanford.edu/~tibs/ElemStatLearn/datasets/zip.info]. It does so by doing a scan that converts the digit (in the form of an image) to grayscale values on a grid. For example using a 16x16 grid, each image would be represented by 256 variables (the intensity of each pixel). For machine-produced digits, this intensity would be quite uniform for each image and the pattern for each digit would be distinct, enabling accurate automatic reading of the digit. Handwritten digits tend to be quite variable (see below). In order to get automatic recognition of handwritten digits, the handwritten digits were scanned and converted to grayscale values on a 16 × 16 grid. The goal is to determine characteristics associated with each digit in order to identify the handwritten digit correctly.

We read the data for all 10 digits into a set of *number of cases* x 256 matrices in such a way that the greyscale values for the first digit are placed in the first row, the second digit in the second row, etc.

```
d.file <- {}
d.digits <- c({})
for (i in 0:9) {
  d.file[i+1] <- paste(data.dir, "/train_", i, ".dat", sep = "")
  d.digits[[i+1]] <- matrix(scan(d.file[i+1], sep = ","), ncol = 256,
                           byrow = T)
}
```

The number of cases varies

```
(num.cases <- unlist(lapply(d.digits, dim))[seq(1,20,2)])
```

## Data Science

```
[1] 1194 1005 731 658 652 556 664 645 542 644
```

Consider the first 144 cases of a few handwritten digits in the dataset.

The `layout` command allows us to create a matrix of images (in this case 12 x 12) and the `byrow = TRUE` indicates that the first 12 images go in the first row, the second 12 in the second row, etc.

The `par(mar = c(0,0,0,0))` command specifies that there will be no margins around the images.

The `matrix(digits[i,],16,16)[,16:1]` command takes each row of `digits`, places them in a 16 x 16 matrix, and re-orders the columns with the `[,16:1]` expression.

The `image` command plots a matrix of values. `col = gray((255:0)/255)` determines the 'blackness' - try different values. See `?image`.)

```
plot.digits <- function(digits) {  
  x11(width = 6, height = 6)      # Open a graphics window of given size  
  # Create a plot matrix with 144 subplots - plot in row-wise order  
  layout(matrix(1:144, 12, 12, byrow = TRUE))  
  # No margin (see ?par)  
  oldpar <- par(mar = c(0,0,0,0))  
  for (i in 1:144) {  
    # xaxt = "n", yaxt = "n" - no axes  
    image(matrix(digits[i,],16,16)[,16:1], xaxt = "n", yaxt = "n",  
             col = gray((255:0)/255))  
  }  
  par(oldpar)  
}
```

We now plot the representations of the digits '2', '3', '5', '8' (the `+1` appears because the digit 0 is in `d.digits[1]`).

```
plot.digits(d.digits[[2+1]])  
plot.digits(d.digits[[3+1]])  
plot.digits(d.digits[[5+1]])  
plot.digits(d.digits[[8+1]])
```




Figure 11.




Figure 12.




Figure 13.


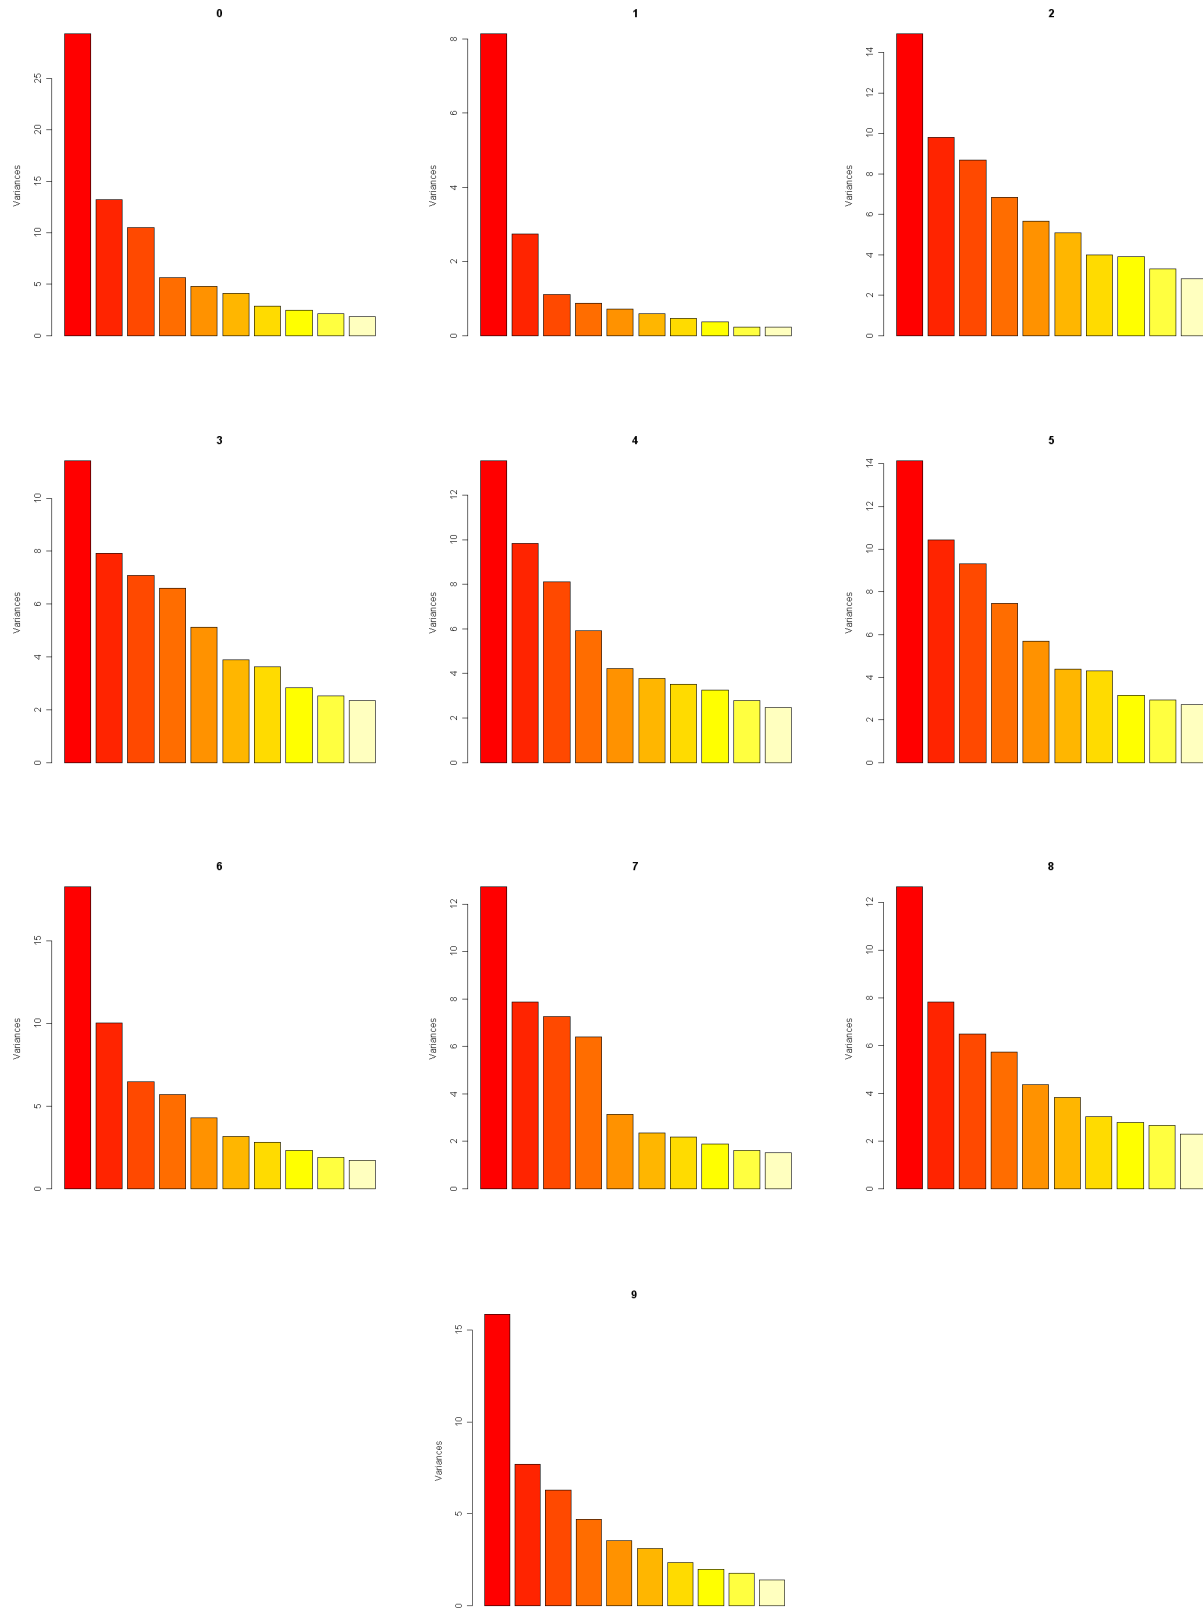


Figure 14.

It is obvious that there is a lot of variation in the shape and thickness.

```
library(stats)
graphics.off()
pc.digits <- {}
for (i in 0:9) {
  pc.digits[[i+1]] <- prcomp(d.digits[[i+1]])
  plot(pc.digits[[i+1]], col = heat.colors(10), main = i)
  print(summary(pc.digits[[i+1]]))
  readline("Press return..")
}
```



**Figure 16.**Principal Components for 0 - 9

This shows the first 10 principal components. Let us look at the principal components as



displayed for the digit 3 by `summary(pc.digits[[i+1]])`.

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
Standard deviation	3.379	2.816	2.6599	2.5680	2.2613	1.9724	1.9066	1.6835
Proportion of Variance	0.127	0.088	0.0785	0.0732	0.0567	0.0432	0.0403	0.0314
Cumulative Proportion	0.127	0.215	0.2931	0.3663	0.4230	0.4662	0.5065	0.5379
	PC9	PC10	PC11	PC12	PC13	PC14	PC15	PC16
Standard deviation	1.5851	1.5287	1.4235	1.3358	1.3052	1.2407	1.1724	1.0906
Proportion of Variance	0.0279	0.0259	0.0225	0.0198	0.0189	0.0171	0.0153	0.0132
Cumulative Proportion	0.5658	0.5917	0.6142	0.6340	0.6529	0.6700	0.6852	0.6984
	PC17	PC18	PC19	PC20	PC21	PC22	PC23	
Standard deviation	1.0706	1.0067	0.9864	0.94142	0.89733	0.89487	0.86294	
Proportion of Variance	0.0127	0.0112	0.0108	0.00983	0.00893	0.00888	0.00826	
Cumulative Proportion	0.7111	0.7224	0.7331	0.74297	0.75190	0.76078	0.76904	
	PC24	PC25	PC26	PC27	PC28	PC29	PC30	
Standard deviation	0.8388	0.82166	0.80241	0.77612	0.75564	0.73613	0.73355	
Proportion of Variance	0.0078	0.00749	0.00714	0.00668	0.00633	0.00601	0.00597	
Cumulative Proportion	0.7769	0.78434	0.79148	0.79816	0.80449	0.81051	0.81647	
	PC31	PC32	PC33	PC34	PC35	PC36	PC37	
Standard deviation	0.72223	0.69970	0.69593	0.67441	0.65674	0.64538	0.6370	
Proportion of Variance	0.00579	0.00543	0.00537	0.00505	0.00478	0.00462	0.0045	
Cumulative Proportion	0.82226	0.82769	0.83306	0.83811	0.84289	0.84751	0.8520	
	PC38	PC39	PC40	PC41	PC42	PC43	PC44	
Standard deviation	0.6225	0.6151	0.59972	0.57805	0.57474	0.57171	0.55860	
Proportion of Variance	0.0043	0.0042	0.00399	0.00371	0.00366	0.00363	0.00346	
Cumulative Proportion	0.8563	0.8605	0.86450	0.86821	0.87187	0.87550	0.87896	
	PC45	PC46	PC47	PC48	PC49	PC50	PC51	
Standard deviation	0.54688	0.53869	0.53151	0.52598	0.51026	0.4936	0.49021	
Proportion of Variance	0.00332	0.00322	0.00313	0.00307	0.00289	0.0027	0.00267	
Cumulative Proportion	0.88228	0.88549	0.88863	0.89170	0.89458	0.8973	0.89995	
	PC52	PC53	PC54	PC55	PC56	PC57	PC58	
Standard deviation	0.47660	0.46972	0.45966	0.45415	0.45177	0.44200	0.43255	
Proportion of Variance	0.00252	0.00245	0.00234	0.00229	0.00226	0.00217	0.00208	
Cumulative Proportion	0.90247	0.90492	0.90726	0.90955	0.91182	0.91398	0.91606	
	PC59	PC60	PC61	PC62	PC63	PC64	PC65	
Standard deviation	0.42818	0.42049	0.4141	0.40813	0.40554	0.40185	0.39848	
Proportion of Variance	0.00203	0.00196	0.0019	0.00185	0.00182	0.00179	0.00176	
Cumulative Proportion	0.91809	0.92005	0.9220	0.92380	0.92563	0.92742	0.92918	
	PC66	PC67	PC68	PC69	PC70	PC71	PC72	
Standard deviation	0.3914	0.38427	0.37629	0.37553	0.36578	0.36246	0.35829	
Proportion of Variance	0.0017	0.00164	0.00157	0.00156	0.00148	0.00146	0.00142	
Cumulative Proportion	0.9309	0.93252	0.93409	0.93565	0.93714	0.93859	0.94002	
	PC73	PC74	PC75	PC76	PC77	PC78	PC79	
Standard deviation	0.3547	0.34802	0.34547	0.34009	0.33361	0.33034	0.32361	
Proportion of Variance	0.0014	0.00134	0.00132	0.00128	0.00123	0.00121	0.00116	
Cumulative Proportion	0.9414	0.94276	0.94408	0.94536	0.94660	0.94781	0.94897	
	PC80	PC81	PC82	PC83	PC84	PC85	PC86	
Standard deviation	0.31983	0.31583	0.31248	0.30989	0.30383	0.30114	0.29579	
Proportion of Variance	0.00113	0.00111	0.00108	0.00107	0.00102	0.00101	0.00097	
Cumulative Proportion	0.95010	0.95121	0.95229	0.95336	0.95438	0.95539	0.95636	
	PC87	PC88	PC89	PC90	PC91	PC92	PC93	
Standard deviation	0.29536	0.29124	0.28636	0.28608	0.28391	0.28096	0.27791	
Proportion of Variance	0.00097	0.00094	0.00091	0.00091	0.00089	0.00088	0.00086	
Cumulative Proportion	0.95733	0.95827	0.95918	0.96009	0.96098	0.96186	0.96271	
	PC94	PC95	PC96	PC97	PC98	PC99	PC100	
Standard deviation	0.27364	0.27051	0.2679	0.26156	0.26002	0.25703	0.25466	
Proportion of Variance	0.00083	0.00081	0.0008	0.00076	0.00075	0.00073	0.00072	
Cumulative Proportion	0.96354	0.96435	0.9651	0.96591	0.96666	0.96739	0.96811	
	PC101	PC102	PC103	PC104	PC105	PC106	PC107	
Standard deviation	0.2515	0.24831	0.24498	0.24344	0.24249	0.23823	0.23705	
Proportion of Variance	0.0007	0.00068	0.00067	0.00066	0.00065	0.00063	0.00062	
Cumulative Proportion	0.9688	0.96950	0.97016	0.97082	0.97147	0.97210	0.97272	
	PC108	PC109	PC110	PC111	PC112	PC113	PC114	

## Data Science

```
Standard deviation    0.23416 0.23096 0.23025 0.22562 0.22385 0.22264 0.22033
Proportion of Variance 0.00061 0.00059 0.00059 0.00056 0.00056 0.00055 0.00054
Cumulative Proportion 0.97333 0.97392 0.97451 0.97508 0.97563 0.97618 0.97672
                        PC115  PC116  PC117  PC118  PC119  PC120  PC121
Standard deviation    0.21758 0.21686 0.21599 0.2115 0.21029 0.20830 0.20638
Proportion of Variance 0.00053 0.00052 0.00052 0.0005 0.00049 0.00048 0.00047
Cumulative Proportion 0.97725 0.97777 0.97829 0.9788 0.97927 0.97975 0.98023
                        PC122  PC123  PC124  PC125  PC126  PC127  PC128
Standard deviation    0.20357 0.20118 0.20011 0.19753 0.19597 0.19338 0.19273
Proportion of Variance 0.00046 0.00045 0.00044 0.00043 0.00043 0.00041 0.00041
Cumulative Proportion 0.98069 0.98114 0.98158 0.98201 0.98244 0.98285 0.98327
                        PC129  PC130  PC131  PC132  PC133  PC134  PC135
Standard deviation    0.19204 0.1903 0.18540 0.18348 0.18224 0.17902 0.17751
Proportion of Variance 0.00041 0.0004 0.00038 0.00037 0.00037 0.00036 0.00035
Cumulative Proportion 0.98367 0.9841 0.98446 0.98483 0.98520 0.98555 0.98590
```

There do not seem to be any dominant components in this case. We see that when using 16 out of 256 components, about 70% of the variance is accounted for; using 51 out of 256 components, about 90% of the variance is accounted for. It takes another 29 components to account for 95% of the variance. We have to determine what our tolerance is (i.e. what proportion of the variance are we wishing to account for?). Note that this produces a lot of output.

```
pc.digits[[3+1]]
```

```
$sdev
```

```
[1] 3.379209e+00 2.816344e+00 2.659944e+00 2.567952e+00 2.261306e+00 1.972429e+00
[7] 1.906591e+00 1.683494e+00 1.585108e+00 1.528654e+00 1.423500e+00 1.335809e+00
[13] 1.305222e+00 1.240664e+00 1.172391e+00 1.090591e+00 1.070565e+00 1.006657e+00
[19] 9.864085e-01 9.414192e-01 8.973307e-01 8.948743e-01 8.629366e-01 8.387566e-01
.
.
[229] 4.935517e-02 4.835835e-02 4.590472e-02 4.424380e-02 4.238942e-02 4.195997e-02
[235] 3.960521e-02 3.839665e-02 3.744329e-02 3.493125e-02 3.469602e-02 3.187392e-02
[241] 2.977113e-02 2.900486e-02 2.816116e-02 2.732529e-02 2.584797e-02 2.445102e-02
[247] 2.096828e-02 1.811004e-02 1.641156e-02 1.448598e-02 1.326148e-02 1.206861e-02
[253] 1.118102e-02 8.358757e-03 1.419367e-03 4.476719e-05
```

```
$rotation
```

```
                PC1                PC2                PC3                PC4                PC5
[1,] 1.099612e-03 -2.819665e-04 7.083808e-04 -6.411722e-04 2.205328e-03
[2,] 7.604992e-03 -1.695856e-03 2.534928e-03 -3.823239e-03 1.083701e-02
[3,] 1.790017e-02 -2.676772e-03 1.286969e-02 6.994898e-03 2.765253e-02
[4,] 2.299110e-02 -1.416062e-02 4.563322e-02 2.643396e-02 5.581772e-02
[5,] 1.034658e-02 -2.710035e-02 7.873074e-02 5.901414e-02 7.673893e-02
[6,] -5.753018e-03 -2.931406e-02 9.973891e-02 8.669514e-02 6.975713e-02
[7,] -1.776670e-02 -2.315438e-02 1.035462e-01 1.071173e-01 4.110756e-02
[8,] -1.759542e-02 -6.651151e-03 9.411738e-02 1.235880e-01 7.910551e-03
[9,] -5.205410e-03 -5.212135e-03 7.991397e-02 1.219229e-01 4.933417e-03
[10,] 2.101444e-02 -7.289502e-03 6.769933e-02 1.007456e-01 2.155735e-02
[11,] 5.211035e-02 -1.133868e-02 5.741599e-02 6.751857e-02 5.483082e-02
[12,] 4.344731e-02 -4.299763e-03 3.059531e-02 2.396153e-02 5.736097e-02
[13,] 2.458593e-02 1.260481e-03 1.308472e-02 -4.973076e-04 3.145985e-02
[14,] 5.425093e-03 -6.667111e-04 3.317569e-03 -3.907695e-03 9.428703e-03
[15,] 3.360513e-04 -3.769016e-04 9.255188e-04 -1.014723e-04 6.460084e-04
[16,] 1.308072e-06 -2.441618e-06 9.819137e-07 6.563653e-07 -1.445957e-06
[17,] 6.610023e-03 -2.103299e-03 3.189284e-03 -3.323865e-03 1.209708e-02
[18,] 2.211880e-02 -7.013420e-03 1.557243e-02 -1.093116e-02 3.541177e-02
[19,] 4.971525e-02 -2.054937e-02 5.344012e-02 1.591546e-02 1.005163e-01
[20,] 6.694901e-02 -6.641992e-02 1.026674e-01 5.569076e-02 1.585556e-01
[21,] 4.391158e-02 -1.014963e-01 1.027717e-01 8.177299e-02 1.499565e-01
[22,] 2.620040e-02 -9.880501e-02 7.411042e-02 7.926029e-02 7.626994e-02
[23,] 2.015773e-02 -8.254495e-02 3.866674e-02 6.348251e-02 2.034777e-02
```

```

[24,] 2.065485e-02 -7.698070e-02 1.547307e-02 5.430896e-02 -6.722872e-03
[25,] 2.293019e-02 -6.868419e-02 1.094210e-02 5.219952e-02 -1.548793e-02
[26,] 3.210304e-02 -4.605961e-02 2.857134e-02 6.772131e-02 -9.114666e-03
[27,] 5.619332e-02 -3.161787e-02 5.490477e-02 9.915164e-02 2.923214e-02
[28,] 9.646655e-02 -2.222523e-02 6.449774e-02 9.504049e-02 9.393785e-02
[29,] 8.776529e-02 -4.057568e-03 3.916974e-02 3.550419e-02 9.051134e-02
[30,] 4.756880e-02 3.595471e-03 1.725299e-02 -7.793875e-03 5.752924e-02
.
.
.
[250,] 9.107206e-02 4.892861e-02 -4.105314e-02 1.400713e-01 -1.583591e-01
[251,] 1.123120e-01 5.401507e-02 -4.764714e-02 8.534363e-02 -1.357300e-01
[252,] 9.436383e-02 3.232665e-02 -4.275065e-02 2.787946e-02 -7.425245e-02
[253,] 5.269707e-02 2.202144e-02 -2.660328e-02 -5.700130e-03 -2.265225e-02
[254,] 2.196619e-02 1.303770e-02 -1.066433e-02 -8.679210e-03 -4.212086e-03
[255,] 4.266127e-03 5.026219e-03 -3.766860e-03 -3.290478e-03 -8.020112e-04
[256,] 5.289045e-05 6.962248e-05 -9.195723e-05 -4.593701e-05 5.537029e-06

```

It may be difficult to understand what these linear combinations mean. So let us look at the first four principal components for each number.

```

pc <- array(dim = c(4, 256, 10), dimnames = list(c(1:4), 1:256, c(0:9)))
for (i in 0:9) {
  pc[1,,i+1] <- pc.digits[[i+1]]$rotation[,1]
  pc[2,,i+1] <- pc.digits[[i+1]]$rotation[,2]
  pc[3,,i+1] <- pc.digits[[i+1]]$rotation[,3]
  pc[4,,i+1] <- pc.digits[[i+1]]$rotation[,4]
}


```

We can take a look at the *average* of all the data along with the *first four principal components* (in this case, the principal component vectors are themselves ‘characters’ of a sort).

```

graphics.off()
x11(width = 4, height = 2.4)
layout(matrix(1:50, 5, 10, byrow = TRUE))
oldpar <- par(mar = c(0,0,0,0))
for (i in 0:9) {
  mean <- apply(d.digits[[i+1]], 2, mean)
  image(matrix(mean,16,16)[,16:1], xaxt = "n", yaxt = "n",
    col = gray((255:0)/255))
  image(matrix(pc[1,,i+1],16,16)[,16:1], xaxt = "n", yaxt = "n",
    col = gray((255:0)/255))
  image(matrix(pc[2,,i+1],16,16)[,16:1], xaxt = "n", yaxt = "n",
    col = gray((255:0)/255))
  image(matrix(pc[3,,i+1],16,16)[,16:1], xaxt = "n", yaxt = "n",
    col = gray((255:0)/255))
  image(matrix(pc[4,,i+1],16,16)[,16:1], xaxt = "n", yaxt = "n",
    col = gray((255:0)/255))
}
par(oldpar)

```



**Figure 17.** Mean and first 4 principal components

The first cell is the mean, the second is  $PC1$ , the third is  $PC2$ , the fourth is  $PC3$  and the fifth is  $PC4$ .

We can look at what happens if we evaluate  $\text{mean} + \lambda * PC1$  ( $-7 \leq \lambda \leq 7$ ).

Because we wish to do this several times, we will create a function to display the mean + pcs of one number.

```
display.mean.pc <- function(pc, digits) {
  mean <- apply(digits, 2, mean)
  for (i in 1:15) {
    image(matrix(mean+(i-8)*pc, 16,16)[,16:1], xaxt = "n",
             yaxt = "n", col = gray((255:0)/255))
  }
}
```

We will also use a function that will put all the numbers using one principal component in the same plot

```
display.pcs <- function (pcnum) {
  x11(width = 7, height = 5)
  oldpar <- par(mar = c(1,0,0,0))
  layout(matrix(1:150, 10, 15, byrow = TRUE))
  for (i in 0:9) {
    display.mean.pc(pc[pcnum,,i+1], d.digits[[i+1]])
  }
  bringToTop(which = dev.cur())
  par(oldpar)
}
display.pcs(1)
```




Figure 18.

It appears that  $PC1$  is generally associated with the *width of the character*. although for '2' it appears to be the height and for '5' the relative widths of the upper and lower halves.

## Data Science

And what happens if we evaluate  $\text{mean} + \mu * PC2$  ( $-7 \leq \mu \leq 7$ )?

`display.pcs(2)`




Figure 19.

In several cases, it appears that  $PC2$  is associated with the *thickness of the character* although in '1', it seems to be the direction of the curve; in '5', the width; in '9', the slope. For several of the digits, it appears to be associated with variability ('2', '4', '5', '6').



Now what happens if we evaluate  $\text{mean} + \mu * PC3$  ( $-7 \leq \mu \leq 7$ )?

`display.pcs(3)`




Figure 20.

It appears that  $PC3$  is associated with the *size of a 'loop' of the character*.

## Data Science

And what happens if we evaluate  $\text{mean} + \mu * PC4$  ( $-7 \leq \mu \leq 7$ )?

```
display.pcs(4)
```




Figure 21.

It appears that for '3', '5', '8'  $PC4$  is associated with the *the relative widths of the upper and lower halves*. For '6' and '9' it looks like total width but it might be relative width.

We can reconstruct our original data (as we noted earlier) using all the principal components, but instead of a full reconstruction, suppose we use a subset of the principal component vectors - for example the first 20. Our first step will be to represent all cases in terms of the new coordinate system (we did this earlier with `pc.1` etc.).

```
d.digits.pc <- {}  
for (i in 0:9) {  
  d.digits.pc[[i+1]] <- d.digits[[i+1]]%*%pc.digits[[i+1]]$rotation  
}
```

For purposes of comparison, we plot the original first 144 images in our dataset (Figure 18).

```
plot.digits(d.digits)
```



Next we find the reconstruction of these images in terms of the first 20 principal components. We use a function to simplify the reconstruction.


We create a new array `tmp` to hold the data. The `cbind(d.digits.pc[[digit+1]][,j])` creates a  $number\ of\ cases \times 1$  **matrix** from the **vector** representing the  $j^{th}$  principal component of our data. This is necessary in order to be able to do the matrix multiplication.) Each component is rotated by multiplying by the  $1 \times 256$  PC vector `pc.digits[[digit+1]]$rotation[,j]` to give a  $number\ of\ cases \times 256$  array, representing the full dataset in the original space. The result for each PC is added to the accumulated results for the previous PCs in `tmp`.

```
recreate <- function(pc.range, digit) {
  tmp <- matrix(0, num.cases[digit+1], 256)
  tmp <-
  d.digits.pc[[digit+1]][,pc.range] %*% t(pc.digits[[digit+1]]$rotation[,pc.range])
  tmp <- tmp/max(abs(range(tmp))) # We want to scale the data to lie in [-1, 1]
  tmp
}
```


The following will use the previous function to reconstruct the images and plot the first 144 of the reconstructed images for the requested digit

- ■ in original form;
- recreated from using only the first 20 principal components;
- recreated from using only the first 100 principal components;
- difference between a 100 PC and a 20 PC reconstruction;
- difference between a 256 PC and a 100 PC reconstruction


```
display.recreate <- function(digit) {
  plot.digits(d.digits[[digit+1]])
  pc.1.20 <- recreate(1:20, digit)
  plot.digits(pc.1.20)
  pc.1.100 <- recreate(1:100, digit)
  plot.digits(pc.1.100)
  pc.21.100 <- recreate(21:100, digit)
  plot.digits(pc.21.100)
  pc.101.256 <- recreate(101:256, digit)
  plot.digits(pc.101.256)
}
```



**Figure 22.** Original data for 1




**Figure 23.** Reconstruction -  
20 PC for 1




**Figure 24.** Reconstruction -  
100 PC for 1

Because '1' is such a simple digit, it is difficult to draw meaningful conclusions from these reconstructions.




**Figure 25.** Difference between a 100 PC  
and a 20 PC reconstruction for '1'.




**Figure 26.** Difference between a 256 PC  
and a 100 PC reconstruction for '1'.


```
display.recreate(3)
```



**Figure 27.** Original data for '3'




**Figure 28.** Reconstruction -  
20 PC for '3'




**Figure 29.** Reconstruction -  
100 PC for '3'

We can see that much of the shape of the characters has been captured by using the first 20 principal components.




**Figure 30.** Difference between a 100 PC  
and a 20 PC reconstruction for '3'.




**Figure 31.** Difference between a 256 PC  
and a 100 PC reconstruction for '3'.

There seems to be little difference between the 100 PC reconstruction and the 256 PC (i.e. complete) reconstruction. Keep in mind that the images have been scaled to enhance the detail as much as possible.




**Figure 32.** Original data for '8'




**Figure 33.** Reconstruction -  
20 PC for '8'



**Figure 34.** Reconstruction -  
100 PC for '8'



**Figure 35.** Difference between a 100 PC  
and a 20 PC reconstruction for '8'.



**Figure 36.** Difference between a 256 PC  
and a 100 PC reconstruction for '8'.

## 4.3 Multidimensional Scaling

Multidimensional scaling (*MDS*) maps data points in  $R^p$  to a lower-dimensional manifold. Consider observations  $x_1, \dots, x_n \in R^p$ ; let  $d_{ij}$  be the distance between observations  $i$  and  $j$  (e.g. Euclidean distance  $d_{ij} = \|x_i - x_j\|$ ). Actually *MDS* needs only some **dissimilarity measure**  $d_{ij}$  between  $x_i$  and  $x_j$ ; it does not need the  $x_i$  and  $x_j$ .

(Other methods such as self-organizing maps (*SOM*), which is related to neural networks, and principal curves and surfaces, which are an extension of principal components, need the actual data points.)

### Kruskal-Shepard Scaling

We seek  $z_1, \dots, z_N \in R^k$  ( $k < p$ ) to minimize *stress* function

$$S_D(z_1, \dots, z_N) = \left[ \sum_{i \neq j} (d_{ij} - \|z_i - z_j\|)^2 \right]^{1/2}.$$

This is *least squares* or *Kruskal-Shepard* scaling. We try to find a lower-dimensional approximation of the data that *preserves the pairwise distances* as much as possible. Note that the approximation is in terms of the distances rather than the squared distances. A gradient descent algorithm is used to minimize  $S_D$ .

### Classical Scaling

In *classical scaling*, (the `cmdscale` in `stats`) we use **similarities**  $s_{ij}$

```
require(stats)
require(MASS)
```

To get an idea of the concept, consider what happens with *projection of a pyramid* versus *MDS on the pyramid*:

- with a projection, the apex would project to the centre and the base corners would remain fixed.
- with MDS, the apex would still project to the centre but the corners may move in order to *try to preserve the relationship of the slant distance to the apex to the base distances* (the higher the apex, the more the corners need to move).

```
(test <- matrix(c(1,1,0, 1,-1,0, -1,1,0, -1,-1,0, 0,0,1), ncol = 3, byrow = T))
      [,1] [,2] [,3]
[1,]    1    1    0
[2,]    1   -1    0
[3,]   -1    1    0
[4,]   -1   -1    0
```



## Data Science

```
[5,] 0 0 1
```

We now find the distances between all pairs of points ( $d_{ij}$ ).

```
(test.dist <- dist(test))
      1      2      3      4
2 2.000000
3 2.000000 2.828427
4 2.828427 2.000000 2.000000
5 1.732051 1.732051 1.732051 1.732051
```

The function `dist`, which is in the `stats` library, produces a lower triangular matrix (with no diagonal elements) which gives the Euclidean distance between every case in the data set.

We see that the distance between adjacent corners is 2, between opposite corners is 2.828427 ( $= 2\sqrt{2}$ ), and from the apex to the corners is 1.732051 ( $= \sqrt{3}$ ).

If we project this pyramid onto the plane, we get a square with a point at the centre.

If we use the classical scaling, we get

```
(test.mds <- cmdscale(test.dist))
      [,1] [,2]
[1,] 0.000000e+00 -1.414214e+00
[2,] -1.414214e+00 -1.372600e-16
[3,] 1.414214e+00 -8.318541e-17
[4,] 1.523304e-16 1.414214e+00
[5,] -2.193883e-16 2.999852e-17
```


These are the  $z$  values.

The distances for the scaled points on the plane are

```
      1      2      3      4
2 2.000000
3 2.000000 2.828427
4 2.828427 2.000000 2.000000
5 1.414214 1.414214 1.414214 1.414214
```

For this pyramid, it turns out that (except for a rotation) MDS and projection produce the same result.

```
plot(test[, -3], xlim = c(-1.5, 1.5), ylim = c(-1.5, 1.5), xlab = "", ylab = "",
      type = "n", main = "Classical vs Projection - apex at 1")
text(test[, 1], test[, 2], 1:5, col = "red")
text(test.mds[, 1], test.mds[, 2], 1:5, col = "blue")
```



**Figure 37.** Blue = projected; red = classical

We will look at the stress later on.

This method uses the centered inner product

$$s_{ij} = (x_i - \bar{x}, x_j - \bar{x})$$

and then minimizes  $\sum_{i \neq j} \{s_{ij} - (z_i - \bar{z}_i, z_j - \bar{z}_j)\}^2$  over  $z_1, \dots, z_N \in R^k$  (the scaled values). There is an explicit solution in terms of eigenvectors. **S** is the centered inner product matrix with elements  $(x_i - \bar{x}, x_j - \bar{x})$ .

To see what the classical MDS does, we take the ‘distance’ values and create a matrix of squared distances as follows

```
tmp <- matrix(0, 5, 5)
row(tmp) > col(tmp)      # to show what 'row(tmp) > col(tmp)' does
  [,1] [,2] [,3] [,4] [,5]
[1,] FALSE FALSE FALSE FALSE FALSE
[2,] TRUE FALSE FALSE FALSE FALSE
[3,] TRUE TRUE FALSE FALSE FALSE
[4,] TRUE TRUE TRUE FALSE FALSE
[5,] TRUE TRUE TRUE TRUE FALSE
```

We populate the lower triangle of the matrix with the squared distance

```
tmp[row(tmp) > col(tmp)] <- test.dist^2
tmp
  [,1] [,2] [,3] [,4] [,5]
[1,] 0 0 0 0 0
[2,] 4 0 0 0 0
[3,] 4 8 0 0 0
[4,] 8 4 4 0 0
[5,] 3 3 3 3 0
```

## Data Science

and then copy the lower triangle to the upper triangle.

```
(S <- tmp + t(tmp))
      [,1] [,2] [,3] [,4] [,5]
[1,]    0    4    4    8    3
[2,]    4    0    8    4    3
[3,]    4    8    0    4    3
[4,]    8    4    4    0    3
[5,]    3    3    3    3    0
```

We have to do a double centering. We start by finding the row, column, and grand means,

```
(grand.mean <- mean(S))
[1] 3.52
(col.mean <- apply(S, 2, mean))
[1] 3.8 3.8 3.8 3.8 2.4
(row.mean <- apply(S, 1, mean))
[1] 3.8 3.8 3.8 3.8 2.4
```

Then we subtract the mean of each column from the columns (notice the use of `t(t(S)-col.mean)` to do the subtraction. The matrices are stored by columns so using `S-col.mean` would subtract `col.mean[1]` from `S[1,1]`, `col.mean[2]` from `S[2,1]`, `col.mean[3]` from `S[3,1]`, and so on. The use of the `t(S)` means that we subtract `col.mean[1]` from `S[1,1]`, `col.mean[2]` from `S[1,2]`, `col.mean[3]` from `S[1,3]`, and so on (as we require). The `t(...)` gives us back the correct result.

```
(S <- t(t(S) - col.mean))
      [,1] [,2] [,3] [,4] [,5]
[1,] -3.8  0.2  0.2  4.2  0.6
[2,]  0.2 -3.8  4.2  0.2  0.6
[3,]  0.2  4.2 -3.8  0.2  0.6
[4,]  4.2  0.2  0.2 -3.8  0.6
[5,] -0.8 -0.8 -0.8 -0.8 -2.4
```

and then the mean of each row from the rows.

```
(S <- S - row.mean)
      [,1] [,2] [,3] [,4] [,5]
[1,] -7.6 -3.6 -3.6  0.4 -3.2
[2,] -3.6 -7.6  0.4 -3.6 -3.2
[3,] -3.6  0.4 -7.6 -3.6 -3.2
[4,]  0.4 -3.6 -3.6 -7.6 -3.2
[5,] -3.2 -3.2 -3.2 -3.2 -4.8
```

The process is completed by adding in the grand mean

```
(S <- S + grand.mean)
      [,1] [,2] [,3] [,4] [,5]
[1,] -4.08 -0.08 -0.08  3.92  0.32
[2,] -0.08 -4.08  3.92 -0.08  0.32
[3,] -0.08  3.92 -4.08 -0.08  0.32
[4,]  3.92 -0.08 -0.08 -4.08  0.32
[5,]  0.32  0.32  0.32  0.32 -1.28
```



We can check our result by taking the row and column means

```
apply(S,1,mean)
[1] -8.874737e-17 -8.876363e-17 -8.876363e-17 -8.874737e-17 8.881784e-17
apply(S,2,mean)
[1] -8.874737e-17 -8.876363e-17 -8.876363e-17 -8.874737e-17 8.881784e-17
```

The small values indicate that the matrix is double centred.

We find the  $k$  largest eigenvalues  $\lambda_1 > \dots > \lambda_k$  of  $\mathbf{S}$ ,

```
k <- 2 # Dimension of the target space
eig <- eigen(-S/2, symmetric = T)
eig$values
[1] 4.000000e+00 4.000000e+00 8.000000e-01 -2.080358e-16 -1.434345e-15
```

with associated eigenvectors  $\mathbf{E}_k = (e_1, \dots, e_k)$ .

We have

```
(E <- eig$vectors[,1:k])
      [,1] [,2]
[1,]  6.762109e-01 -2.067338e-01
[2,] -2.067338e-01 -6.762109e-01
[3,]  2.067338e-01  6.762109e-01
[4,] -6.762109e-01  2.067338e-01
[5,] -1.301043e-18  4.580077e-17
```

Let  $\mathbf{D}_k$  be a diagonal matrix with diagonal entries  $\sqrt{\lambda_i}, i = 1, \dots, k$ .

```
(D <- diag(sqrt(eig$values[1:k])))
      [,1] [,2]
[1,]     2     0
[2,]     0     2
```

The solutions  $z_i$  to the classical scaling problem are the rows of  $\mathbf{E}_k \mathbf{D}_k$


```
(cmds <- E%*%D)
      [,1] [,2]
[1,]  1.352422e+00 -4.134676e-01
[2,] -4.134676e-01 -1.352422e+00
[3,]  4.134676e-01  1.352422e+00
[4,] -1.352422e+00  4.134676e-01
[5,] -2.602085e-18  9.160153e-17
```

These are not the same values as we obtained from `cmdscale` (above) but, it is equivalent as can be seen from the fact that the distances are the same (it is very close to the projected values).

```
dist(cmds)
      1      2      3      4
2 2.000000
3 2.000000 2.828427
4 2.828427 2.000000 2.000000
5 1.414214 1.414214 1.414214 1.414214
```

## Data Science

```
plot(test[, -3], xlim = c(-1.5, 1.5), ylim = c(-1.5, 1.5), xlab = "", ylab = "",
      type = "n", main = "cmdscale vs eigen - apex at 1")
text(cmds[, 1], cmds[, 2], 1:5, col = "red")
text(test.mds[, 1], test.mds[, 2], 1:5, col = "blue")
```




**Figure 38.** Blue = projected; red = classical

(Note that *classical scaling is not equivalent to least squares scaling*, since inner products rely on a choice of origin while pairwise distances do not; a set of inner products determines a set of pairwise distances but not vice versa.)

We might wonder if the projection and classical MDS methods will produce the same result in general. To investigate this, move the apex to 2.236. There is a function that can be used to do this.

```
source(paste(code.dir, "ClassicMDS.r", sep = "/"))
Classic.MDS(2.236)
```


When we use the same process as before, we find that we do in fact get the same distance information and a similar image



**Figure 39.** Blue = projected; red = classical

The fun begins if we try again with

`Classic.MDS(2.237)`



**Figure 40.** Blue = projected; red = classical

## Data Science

If we look at the distance as printed by the function

```
[1] "dist(cmds)"
      1          2          3          4
2 1.8201571
3 0.8288717 0.9912854
4 2.6490287 0.8288717 1.8201571
5 2.5997129 2.2912509 2.2912509 2.5997129
```

we no longer have the original distance values (as given above) of

```
      1          2          3          4
2 2.000000
3 2.000000 2.828427
4 2.828427 2.000000 2.000000
5 1.414214 1.414214 1.414214 1.414214
```

but we also find that the distance values for the eigenvalue method are different, as can be seen from the figure. **WHY?**

There are two factors at work here:

- We believe that we should project from three dimensions to two dimensions (otherwise we have no dimension reduction), but the number of dimensions required is determined by the *k* largest eigenvalues. We saw in the original case that the eigenvalues were [4, 4, 0.8, 0, 0] so that the first two dimensions are dominant. What we did not look at in the 2.236 and 2.237 cases were the eigenvalues which were [4, 4, 3.999757, 0, 0] and [4.003335, 4, 4, 0, 0] (look at the printed output). We see that the 2.236 case was on the borderline of having 2 eigenvalues that are larger than the others, while the 2.237 case suggests that two dimensions are no longer adequate.
- The *cmdscale* method and the *eigenvalue* method should produce the same results but are different. A closer look at the process in both cases indicates that the matrices for which the eigenvectors are found differ only by amount in the order of  $10^{-16}$  and yet they produce somewhat different results. ***We need to watch out for numerical instabilities.***

## Sammon Mapping

An alternative to the *Kruskal-Shepard* scaling and the *classical scaling* is the *Sammon mapping* which minimizes


$$\sum_{i \neq j} \frac{(d_{ij} - \|z_i - z_j\|)^2}{d_{ij}}$$

and places more emphasis on **preserving smaller pairwise distances**.

We will apply this method to the vertex at 2.237 case.

```
test <- matrix(c(1,1,0, 1,-1,0, -1,1,0, -1,-1,0, 0,0,2.237), ncol = 3, byrow = T)
test.dist <- dist(test)
(test.mds <- sammon(test.dist))
Initial stress      : 0.11184
stress after  10 iters: 0.06046, magic = 0.500
stress after  20 iters: 0.06037, magic = 0.500
$points
      [,1]      [,2]
[1,] -1.0857805 -8.263581e-01
[2,]  0.1748359 -1.412080e+00
[3,]  0.1748359  1.412080e+00
[4,] -1.0857805  8.263581e-01
[5,]  1.8218892 -2.586420e-13
$stress
[1] 0.06036759
$call
sammon(d = test.dist)
plot(test[, -3], xlim = c(-1.5,1.8), ylim = c(-1.5,1.5), xlab = "", ylab = "",
      type = "n", main = "Projected vs Sammon - apex at 2.237")
text(test[,1],test[,2],1:5, col = "red")
text(test.mds$points[,1],test.mds$points[,2],1:5, col = "blue")
li <- c(1,2,5,4,3,5,1,3,4,2)
lines(test[li,1],test[li,2], col = "red", lwd = 2)
lines(test.mds$points[li,1],test.mds$points[li,2],col = "blue", lwd = 2)
```

Projected vs Sammon - apex at 2.237

**Figure 41.** Blue = projected; red = Sammon

The lines indicate the projected pyramid.

```
(dist.sam <- dist(test.mds$points))
      1      2      3      4
2 1.390045
3 2.569000 2.824160
4 1.652716 2.569000 1.390045
5 3.022815 2.169506 2.169506 3.022815
```

To illustrate what is taking place we note that the `sammon` routine in `MASS` uses the `cmdscale` from `stats` as a starting configuration.

```
(test.mds <- cmdscale(test.dist))
      [,1]      [,2]
[1,] -0.4474 -4.050837e-01
[2,] -0.4474 -1.354957e+00
[3,] -0.4474  1.354957e+00
[4,] -0.4474  4.050837e-01
[5,]  1.7896 -7.862374e-14
```

To find the initial stress, we use

```
(total <- sum(test.dist))
[1] 24.24301
```

Compute  $d_{ij} - \|z_i - z_j\|$ .

```
(diff <- test.dist - dist(cmdscale(test.dist)))
      1      2      3      4
2 1.05012715
3 0.23995978 0.11851406
4 2.01825975 0.23995978 1.05012715
5 0.37315792 0.03118513 0.03118513 0.37315792
```

Next compute

$$\sum_{i \neq j} \frac{(d_{ij} - \|z_i - z_j\|)^2}{d_{ij}}$$

```
(err <- sum((diff^2)/test.dist))
[1] 2.711433
```

The stress function that is used in the Sammon routine is a **scaled stress**,

```
err/total
[1] 0.1118439
```

The routine then uses a gradient descent to minimize the stress.

The final stress is found by

```
(diffs <- test.dist-dist.sam)
      1          2          3          4
2  0.609955357
3 -0.568999721  0.004266701
4  1.175711004 -0.568999721  0.609955357
5 -0.376275909  0.477033374  0.477033374 -0.376275909
(errs <- sum((diffs^2)/test.dist))
[1] 1.463492
errs/total          #stress
[1] 0.06036759
```

The effect of this mapping is to project the 5 points of the 3-dimensional pyramid onto the plane in such a way that the relative differences between the true and scaled distances are as small as possible. Note that the classical routine produced a set of points in the plane that

made the differences between those points in the plane and the scaled points the same, but did not minimize the off-plane distances.

Least squares and classical scaling are *metric* scaling methods ( the actual dissimilarities or similarities are approximated.)

**Shepard-Kruskal nonmetric** scaling effectively uses only *ranks*. Nonmetric scaling minimizes the stress function

$$\sum_{i,j} \frac{[\theta(\|z_i - z_j\|) - d_{ij}]^2}{\sum_{i,j} d_{ij}^2}$$

over the  $d_{ij}$  and an arbitrary increasing function  $\theta(\cdot)$ . Fixing  $\theta(\cdot)$ , we use gradient descent to minimize over  $d_{ij}$ . Fixing  $d_{ij}$ , we use **isotonic regression** to find the best monotonic approximation  $\theta(\cdot)$ . We iterate these steps until the solutions seem to stabilize.

**Note:** In **principal surfaces** and **SOM**, points close together in our original space should map close together in the manifold, but points far apart in the original space *might also map close together*. This is less likely in *MDS* since it explicitly tries to preserve all pairwise distances.



## Examples:

Consider a situation in which you do not (or can not) know the data but do know the dissimilarities.

For example, we might have the following table of distances between European cities as found in the dataset `eurodist`.

```
data(eurodist)
eurodist
```

	Athens	Barcelona	Brussels	Calais	Cherbourg	Cologne	Copenhagen
Barcelona	3313						
Brussels	2963	1318					
Calais	3175	1326	204				
Cherbourg	3339	1294	583	460			
Cologne	2762	1498	206	409	785		
Copenhagen	3276	2218	966	1136	1545	760	
Geneva	2610	803	677	747	853	1662	1418
Gibraltar	4485	1172	2256	2224	2047	2436	3196
Hamburg	2977	2018	597	714	1115	460	460
Hook of Holland	3030	1490	172	330	731	269	269
Lisbon	4532	1305	2084	2052	1827	2290	2971
Lyons	2753	645	690	739	789	714	1458
Madrid	3949	636	1558	1550	1347	1764	2498
Marseilles	2865	521	1011	1059	1101	1035	1778
Milan	2282	1014	925	1077	1209	911	1537
Munich	2179	1365	747	977	1160	583	1104
Paris	3000	1033	285	280	340	465	1176
Rome	817	1460	1511	1662	1794	1497	2050
Stockholm	3927	2868	1616	1786	2196	1403	650
Vienna	1991	1802	1175	1381	1588	937	1455

```
Geneva Gibraltar Hamburg Hook of Holland Lisbon Lyons Madrid
```

	Geneva	Gibraltar	Hamburg	Hook of Holland	Lisbon	Lyons	Madrid
Barcelona							
Brussels							
Calais							
Cherbourg							
Cologne							
Copenhagen							
Geneva							
Gibraltar	1975						
Hamburg	1118	2897					
Hook of Holland	895	2428	550				
Lisbon	1936	676	2671	2280			
Lyons	158	1817	1159	863	1178		
Madrid	1439	698	2198	1730	668	1281	
Marseilles	425	1693	1479	1183	1762	320	1157
Milan	328	2185	1238	1098	2250	328	1724
Munich	591	2565	805	851	2507	724	2010
Paris	513	1971	877	457	1799	471	1273
Rome	995	2631	1751	1683	2700	1048	2097
Stockholm	2068	3886	949	1500	3231	2108	3188
Vienna	1019	2974	1155	1205	2937	1157	2409

```
Marseilles Milan Munich Paris Rome Stockholm
```

	Marseilles	Milan	Munich	Paris	Rome	Stockholm
Barcelona						
Brussels						
Calais						
Cherbourg						
Cologne						
Copenhagen						
Geneva						
Gibraltar						

## Data Science

Hamburg						
Hook of Holland						
Lisbon						
Lyons						
Madrid						
Marseilles						
Milan	618					
Munich	1109	331				
Paris	792	856	821			
Rome	1011	586	946	1476		
Stockholm	2428	2187	1754	1827	2707	
Vienna	1363	898	428	1249	1209	2105

We will use multidimensional scaling on this data. In order to plot the results on a map of Europe, we will need to do some scaling of the results to make them fit on the map.

The following gives us a way of plotting images (in this case Portable GreyMap). We will use the classical, iso, and Sammon mappings.

```
library(pixmap)
d.file <- paste(data.dir, "Europe.pgm", sep = "/")
```

For Figure 42:

```
image <- read.pnm(d.file)
plot(image, main = "Classical MDS")
loc.cmd <- cmdscale(eurodist)
x <- (loc.cmd[,1]-min(loc.cmd[,1]))*.14 + 50
y <- -(loc.cmd[,2]-max(loc.cmd[,2]))*.12 + 80
text(x, y, labels(eurodist), cex = 1, col = "red")
```

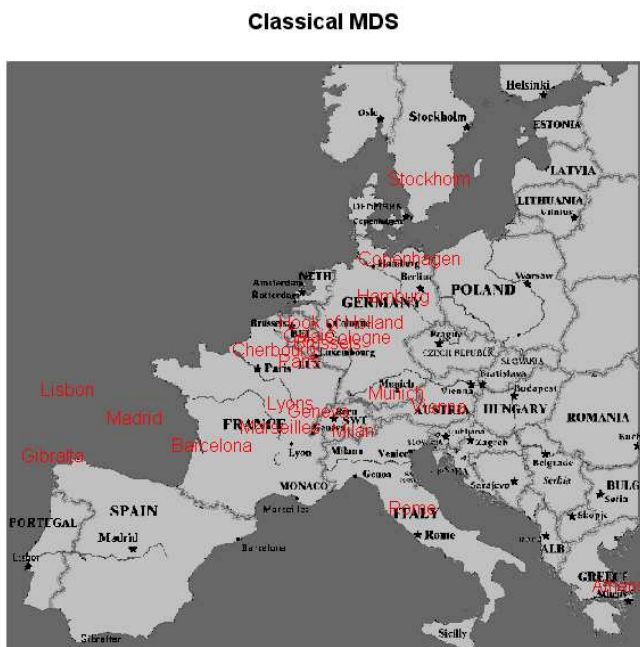


Figure 42.

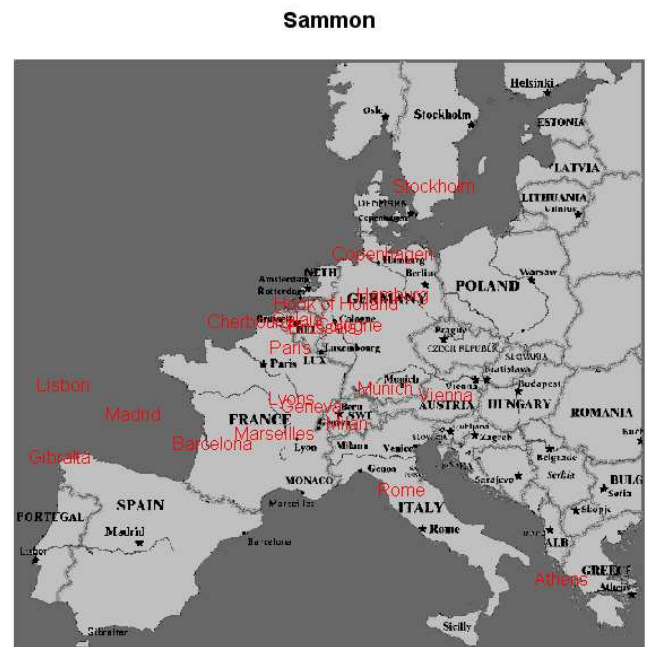


Figure 43.

For Figure 43:

```
plot(image, main = "Sammon")
loc <- sammon(eurodist)
locPt <- loc$points
x <- (loc$points[,1]-min(loc$points[,1]))*.14 + 50
y <- -(loc$points[,2]-max(loc$points[,2]))*.12 + 80
text(x, y, labels(eurodist), cex = 1, col = "red")
```

For Figure 44:

```
plot(image, main = "isoMDS")
loc <- isoMDS(eurodist)
locPt <- loc$points
x <- (locPt[,1]-min(locPt[,1]))*.14 + 50
y <- -(locPt[,2]-max(locPt[,2]))*.15 + 80
text(x, y, labels(eurodist), cex = 1, col = "red")
```

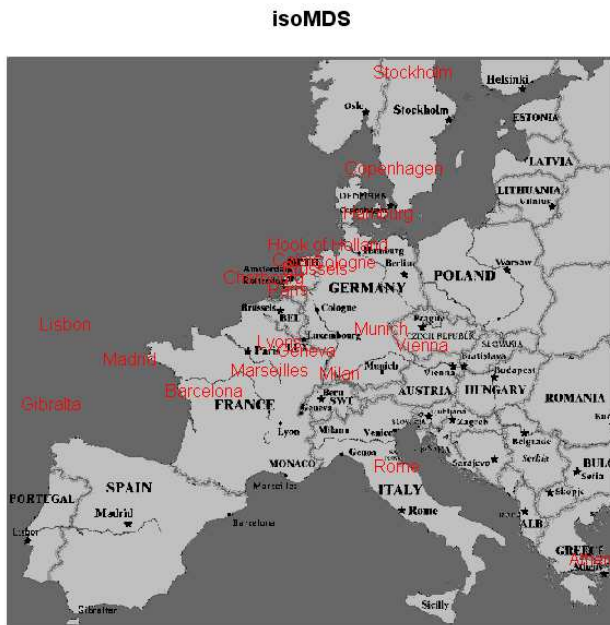


Figure 44.

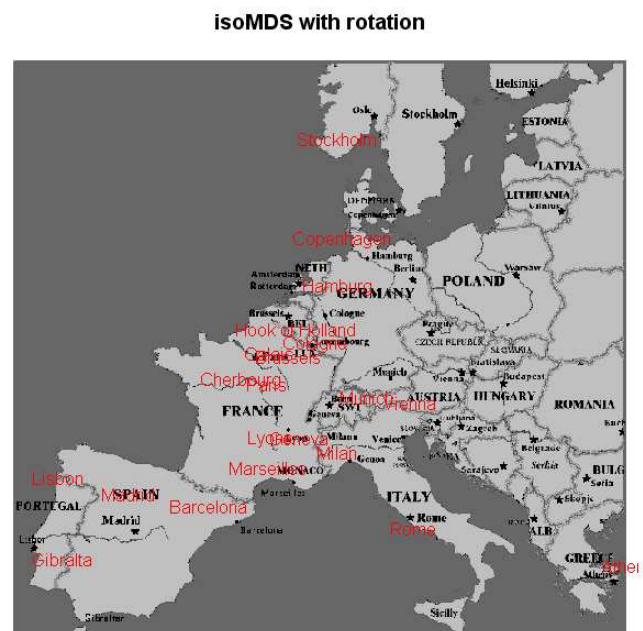


Figure 45.

As can be seen in Figure 44, the information on the distances between points (cities) allows us to place the cities reasonably well. Keep in mind that we are only obtaining relative locations and that the use of rotation might improve the “map”.

For Figure 45:

```
source(paste(code.dir, "3Drotations.r", sep = "/"))
plot(image, main = "isoMDS with rotation")
loc <- isoMDS(eurodist)
locPt <- R.2D(loc$points, -.25)
x <- (locPt[,1]-min(locPt[,1]))*.14 + 50
y <- -(locPt[,2]-max(locPt[,2]))*.15 + 80
text(x, y, names(eurodist), cex = 1, col = "red")
```

### Soft drink taste test example

Consider a taste test in which 10 students did a taste test on 10 soft drinks Diet Pepsi, RC Cola, Yukon, Dr. Pepper, Shasta, Coca-Cola, Diet Dr. Pepper, Tab, Pepsi-Cola, Diet-Rite. The similarity matrix represents the perception of the students as to the similarity of the tastes.

```
comp <- c("Diet Pepsi", "RC Cola", "Yukon", "Dr. Pepper", "Shasta", "Coca-Cola",  
         "Diet Dr. Pepper", "Tab", "Pepsi-Cola", "Diet-Rite")  
p <- length(comp)  
dat <- scan(paste(data.dir, "softdrinks.dat", sep = "/"))  
par(mfrow = c(2,3), mar = c(1,0,1,0), xaxt = "n", yaxt = "n")  
k <- 1  
  # Repeat 9 times to get 9 students  
for (kk in 1:9) {  
  Dis <- matrix(0, p, p)  
  Dis[col(Dis) >= row(Dis)] <- dat[k:(k+54)]  
  Dis <- t(Dis)+Dis  
  diag(Dis) <- diag(Dis/2)  
  k <- k + 55  
  dimnames(Dis) <- list(comp, comp)  
  coords <- cmdscale(Dis)  
  coord1 <- -coords[,1]  
  coord2 <- -coords[,2]  
  plot(coord2, coord1, xlab = "", main = "cmdscale")  
  text(coord2, coord1, comp)  
  coords <- isoMDS(Dis)  
  coord1 <- -coords$points[,1]  
  coord2 <- -coords$points[,2]  
  plot(coord2, coord1, xlab = "", main = paste("isoMDS ", kk))  
  text(coord2, coord1, comp)  
  coords <- sammon(Dis)  
  coord1 <- -coords$points[,1]  
  coord2 <- -coords$points[,2]  
  plot(coord2, coord1, xlab = "", main = "sammon")  
  text(coord2, coord1, comp)  
  readline("Press any key")  
}  
par(oldpar)
```

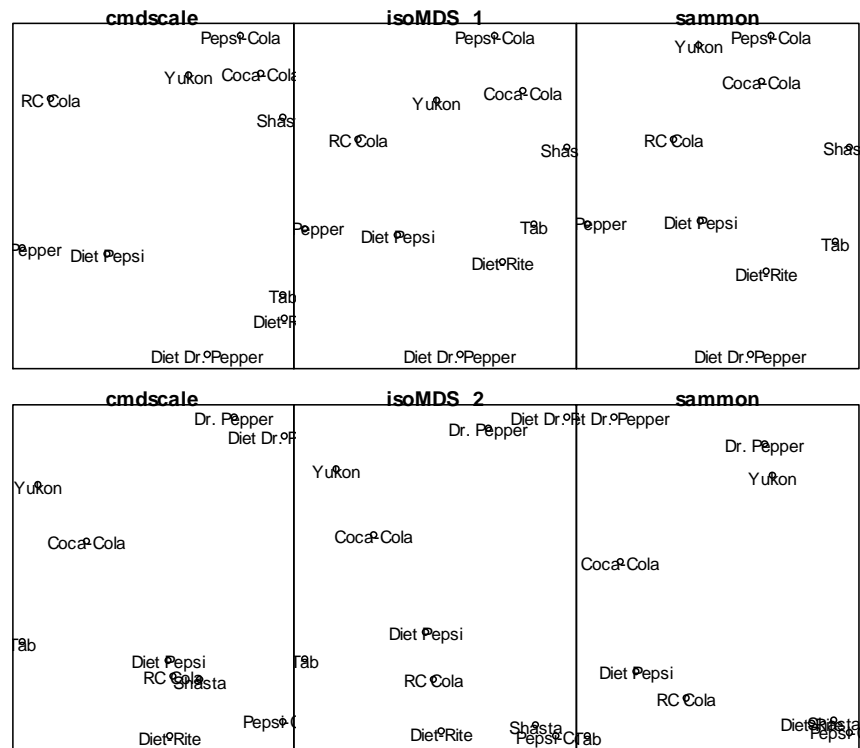


Figure 46.

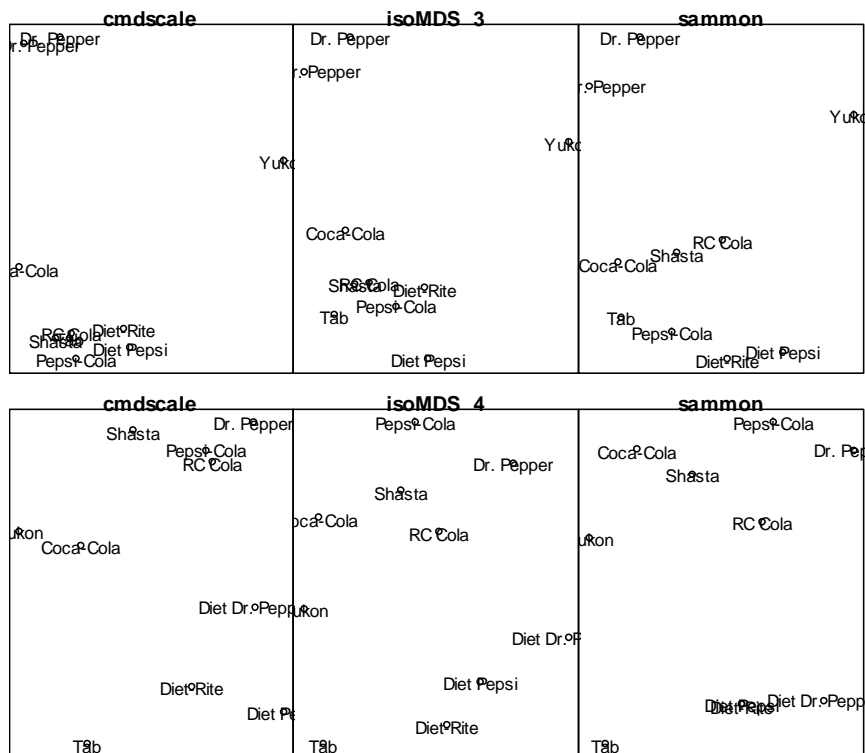


Figure 47.

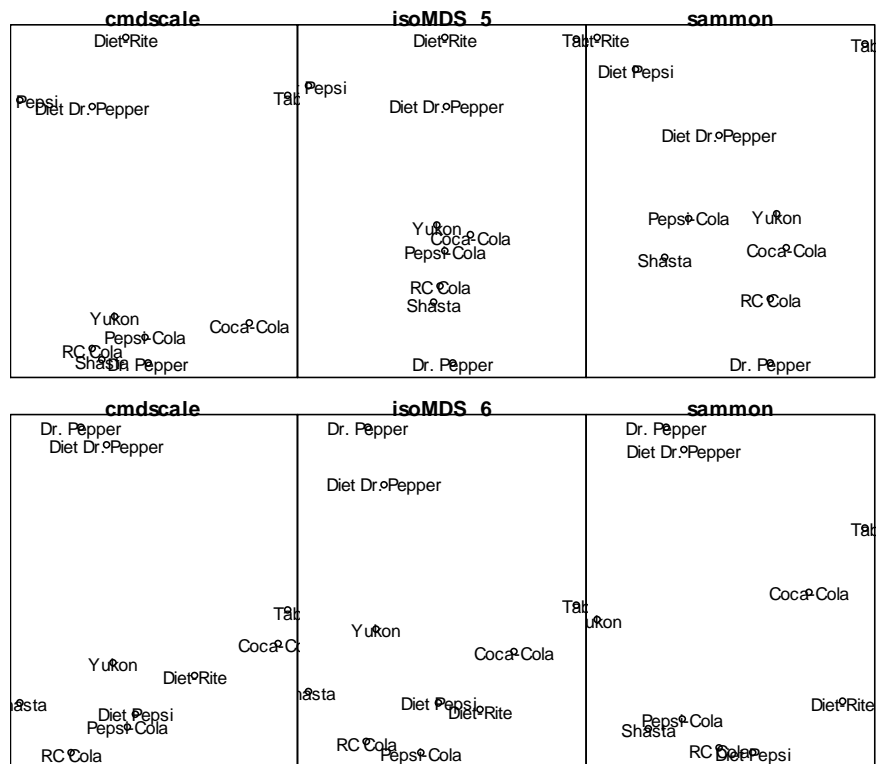


Figure 48.




Figure 49.

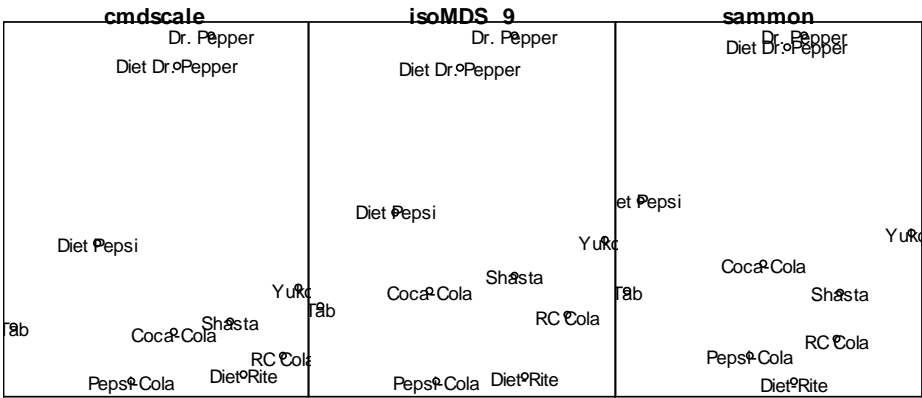


Figure 50.



## Flea Beetles

As a different type of example, we can look at the flea beetle data (again!).

In this case we have high dimensional (6) data that has clustering.



```
source(paste(code.dir, "ReadFleas.r", sep = "/"))
(flea.dist <- dist(d.flea))
```

	1	2	3	4	5	6	7	8
2	8.246211							
3	12.609520	16.031220						
4	20.856654	16.583124	29.35984					
5	24.535688	22.135944	36.18011	16.03122				
6	35.930488	31.288976	44.85532	16.73320	18.62794			
7	7.549834	9.539392	14.89966	18.54724	25.27845	34.813790		
8	10.908712	8.306624	18.05547	18.65476	21.61018	30.724583	15.684387	
9	21.702534	16.278821	30.13304	19.64688	19.20937	26.381812	24.124676	13.638182
10	34.942810	31.352831	43.95452	16.18641	17.97220	6.164414	33.555923	31.112698
11	16.000000	13.490738	12.44990	23.89561	34.46738	39.433488	13.820275	18.411953
12	17.291616	11.958261	27.60435	10.86278	13.89244	22.583180	16.492423	14.966630
13	12.206556	17.291616	11.48913	31.43247	33.00000	45.055521	18.654758	15.874508
14	5.000000	9.433981	15.29706	22.71563	24.63737	37.496667	8.944272	13.190906
15	9.848858	5.196152	19.79899	12.96148	17.52142	27.018512	11.224972	7.615773
16	17.435596	14.560220	19.77372	13.96424	26.83282	28.053520	15.459625	16.401219
17	22.022716	16.941074	27.34959	13.49074	21.97726	21.118712	23.108440	14.000000
...								

```
flea.mds <- isoMDS(flea.dist)
initial value 4.053595
final value 3.626536
converged
```

The `plot` command below sets up the plot but the `type = "n"` prevents any data from being displayed.


```
plot(flea.mds$points, type = "n", main = "isoMDS for Flea Beetles")
text(flea.mds$points, labels = as.character(1:nrow(d.flea)), col = species+1)
```

**Figure 51.****Figure 52.**

```
flea.sam <- sammon(dist(flea.dist))
Initial stress : 0.02439
stress after 9 iters: 0.01203
plot(flea.sam$points, type = "n", main = "Sammon for Flea Beetles")
text(flea.sam$points, labels = as.character(1:nrow(d.flea)), col = species+1)
```

For comparison, we can look at the projection on the plane that produces one of the best separations of the species.

```
plot(-d.flea[,c(1,6)], col = species+1)
```

**Figure 53.**

The use of multidimensional scaling may enable us to see the clusters with better separation.